

KPMG VIRTUAL INTERNSHIP PROJECT

TASK: 1 - Data Quality Assessment

Assessment of data quality and completeness in preparation for analysis.

The client provided KPMG with 3 datasets:

- 1.Customer Demographic
- 2.Customer Addresses
- 3.Transactions data in the past 3 months

```
In [1]: # Importing the required libraries
import pandas as pd
import numpy as np
```

Reading files

```
In [9]: data = pd.ExcelFile("KPMG.xlsx")
```

Reading each file separately

```
In [12]: Transactions = pd.read_excel(data, 'Transactions', header=1)
NewCustomerList = pd.read_excel(data, 'NewCustomerList', header=1)
CustomerDemographic = pd.read_excel(data, 'CustomerDemographic', header=1)
CustomerAddress = pd.read_excel(data, 'CustomerAddress', header=1)
```

C:\Users\Jay\AppData\Local\Temp\ipykernel_4212\2496028931.py:2: FutureWarning: Inferring datetime64[ns] from data containing strings is deprecated and will be removed in a future version. To retain the old behavior explicitly pass Series(data, dtype={value.dtype})

```
NewCustomerList = pd.read_excel(data, 'NewCustomerList', header=1)
```

C:\Users\Jay\AppData\Local\Temp\ipykernel_4212\2496028931.py:3: FutureWarning: Inferring datetime64[ns] from data containing strings is deprecated and will be removed in a future version. To retain the old behavior explicitly pass Series(data, dtype={value.dtype})

```
CustomerDemographic = pd.read_excel(data, 'CustomerDemographic', header=1)
```

Exploring Transactions Data Set

```
In [16]: Transactions.head(5)
```

```
Out[16]:
```

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	price
0	1	2	2950	2017-02-25	0.0	Approved	Solex	

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	pro
1	2	3	3120	2017-05-21	1.0	Approved	Trek Bicycles	
2	3	37	402	2017-10-16	0.0	Approved	OHM Cycles	
3	4	88	3135	2017-08-31	0.0	Approved	Norco Bicycles	
4	5	78	787	2017-10-01	1.0	Approved	Giant Bicycles	

In [17]:

Transactions.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   transaction_id                        20000 non-null  int64
1   product_id                           20000 non-null  int64
2   customer_id                          20000 non-null  int64
3   transaction_date                     20000 non-null  datetime64[ns]
4   online_order                         19640 non-null  float64
5   order_status                         20000 non-null  object
6   brand                                19803 non-null  object
7   product_line                         19803 non-null  object
8   product_class                       19803 non-null  object
9   product_size                        19803 non-null  object
10  list_price                          20000 non-null  float64
11  standard_cost                       19803 non-null  float64
12  product_first_sold_date             19803 non-null  float64
dtypes: datetime64[ns](1), float64(4), int64(3), object(5)
memory usage: 2.0+ MB
```

In [18]:

```
#Using only the required columns
Transactions = Transactions.iloc[:, 0:13]
Transactions.head()
```

Out[18]:

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand	pro
0	1	2	2950	2017-02-25	0.0	Approved	Solex	
1	2	3	3120	2017-05-21	1.0	Approved	Trek Bicycles	
2	3	37	402	2017-10-16	0.0	Approved	OHM Cycles	
3	4	88	3135	2017-08-31	0.0	Approved	Norco Bicycles	
4	5	78	787	2017-10-01	1.0	Approved	Giant Bicycles	

In [19]:

Transactions.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   transaction_id                        20000 non-null  int64
1   product_id                           20000 non-null  int64
2   customer_id                          20000 non-null  int64
3   transaction_date                     20000 non-null  datetime64[ns]
4   online_order                         19640 non-null  float64
5   order_status                         20000 non-null  object
6   brand                               19803 non-null  object
7   product_line                         19803 non-null  object
8   product_class                       19803 non-null  object
9   product_size                         19803 non-null  object
10  list_price                           20000 non-null  float64
11  standard_cost                       19803 non-null  float64
12  product_first_sold_date              19803 non-null  float64
dtypes: datetime64[ns](1), float64(4), int64(3), object(5)
memory usage: 2.0+ MB
```

```
In [20]: #Checking the shape of the data
         Transactions.shape
```

```
Out[20]: (20000, 13)
```

```
In [21]: #Checking for null values
         Transactions.isnull().sum()
```

```
Out[21]: transaction_id      0
         product_id        0
         customer_id       0
         transaction_date   0
         online_order      360
         order_status       0
         brand             197
         product_line       197
         product_class      197
         product_size       197
         list_price         0
         standard_cost      197
         product_first_sold_date 197
         dtype: int64
```

There are missing values in 7 columns. They can be dropped or treated according to the nature of analysis

```
In [22]: #Checking for duplicate values
         Transactions.duplicated().sum()
```

```
Out[22]: 0
```

There are no duplicate values, so the data is unique.

```
In [23]: #check for uniqueness of each column
         Transactions.nunique()
```

```
Out[23]: transaction_id      20000
         product_id         101
```

```

customer_id          3494
transaction_date      364
online_order          2
order_status          2
brand                 6
product_line          4
product_class         3
product_size          3
list_price            296
standard_cost         103
product_first_sold_date 100
dtype: int64

```

Exploring the columns

In [27]: `Transactions.columns`

Out[27]: `Index(['transaction_id', 'product_id', 'customer_id', 'transaction_date', 'online_order', 'order_status', 'brand', 'product_line', 'product_class', 'product_size', 'list_price', 'standard_cost', 'product_first_sold_date'], dtype='object')`

In [28]: `Transactions['order_status'].value_counts()`

Out[28]: `Approved 19821
Cancelled 179
Name: order_status, dtype: int64`

In [29]: `Transactions['brand'].value_counts()`

Out[29]: `Solex 4253
Giant Bicycles 3312
WeareA2B 3295
OHM Cycles 3043
Trek Bicycles 2990
Norco Bicycles 2910
Name: brand, dtype: int64`

In [30]: `Transactions['product_line'].value_counts()`

Out[30]: `Standard 14176
Road 3970
Touring 1234
Mountain 423
Name: product_line, dtype: int64`

In [31]: `Transactions['product_class'].value_counts()`

Out[31]: `medium 13826
high 3013
low 2964
Name: product_class, dtype: int64`

In [32]: `Transactions['product_size'].value_counts()`

Out[32]: `medium 12990
large 3976`

small 2837
 Name: product_size, dtype: int64

In [33]: Transactions['product_first_sold_date']

Out[33]:

0	41245.0
1	41701.0
2	36361.0
3	36145.0
4	42226.0
	...
19995	37823.0
19996	35560.0
19997	40410.0
19998	38216.0
19999	36334.0

Name: product_first_sold_date, Length: 20000, dtype: float64

In [34]: *#convert date column from integer to datetime*
 Transactions['product_first_sold_date'] = pd.to_datetime(Transactions['product_first_sold_date'])
 Transactions['product_first_sold_date'].head()

Out[34]:

0	1970-01-01 11:27:25
1	1970-01-01 11:35:01
2	1970-01-01 10:06:01
3	1970-01-01 10:02:25
4	1970-01-01 11:43:46

Name: product_first_sold_date, dtype: datetime64[ns]

In [35]: Transactions['product_first_sold_date'].head(20)

Out[35]:

0	1970-01-01 11:27:25
1	1970-01-01 11:35:01
2	1970-01-01 10:06:01
3	1970-01-01 10:02:25
4	1970-01-01 11:43:46
5	1970-01-01 10:50:31
6	1970-01-01 09:29:25
7	1970-01-01 11:05:15
8	1970-01-01 09:17:35
9	1970-01-01 10:36:56
10	1970-01-01 11:19:44
11	1970-01-01 11:42:52
12	1970-01-01 09:35:27
13	1970-01-01 09:36:26
14	1970-01-01 10:36:33
15	1970-01-01 10:31:13
16	1970-01-01 10:36:46
17	1970-01-01 09:24:48
18	1970-01-01 11:05:15
19	1970-01-01 10:22:17

Name: product_first_sold_date, dtype: datetime64[ns]

The values in the product_first_sold_date columns are not correct as it shows everything happening the same day at different times.

Exploring New Customer List Data Set

In [36]:

NewCustomerList.head(5)

Out[36]:

	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job_title	job_indu
0	Chickie	Brister	Male	86	1957-07-12	General Manager	
1	Morly	Genery	Male	69	1970-03-22	Structural Engineer	
2	Ardelis	Forrester	Female	10	1974-08-28	Senior Cost Accountant	Fir
3	Lucine	Stutt	Female	64	1979-01-28	Account Representative III	
4	Melinda	Hadlee	Female	34	1965-09-21	Financial Analyst	Fir

5 rows × 23 columns



In [37]:

NewCustomerList.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 23 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   first_name                                                            1000 non-null   object
1   last_name                                                             971 non-null    object
2   gender                                                                1000 non-null   object
3   past_3_years_bike_related_purchases  1000 non-null   int64
4   DOB                                                                  983 non-null    datetime64[ns]
5   job_title                                                             894 non-null    object
6   job_industry_category                                                835 non-null    object
7   wealth_segment                                                       1000 non-null   object
8   deceased_indicator                                                   1000 non-null   object
9   owns_car                                                             1000 non-null   object
10  tenure                                                               1000 non-null   int64
11  address                                                             1000 non-null   object
12  postcode                                                            1000 non-null   int64
13  state                                                               1000 non-null   object
14  country                                                             1000 non-null   object
15  property_valuation                                                    1000 non-null   int64
16  Unnamed: 16                                                           1000 non-null   float64
17  Unnamed: 17                                                           1000 non-null   float64
18  Unnamed: 18                                                           1000 non-null   float64
19  Unnamed: 19                                                           1000 non-null   float64
20  Unnamed: 20                                                           1000 non-null   int64
21  Rank                                                                1000 non-null   int64
22  Value                                                                1000 non-null   float64
dtypes: datetime64[ns](1), float64(5), int64(6), object(11)
memory usage: 179.8+ KB
```

In [38]:

```
#Dropping the unnamed columns
NewCustomerList.drop(['Unnamed: 16', 'Unnamed: 17', 'Unnamed: 18',
                      'Unnamed: 19', 'Unnamed: 20'], axis=1, inplace=True)
```

```
In [39]: #Checking the shape of the dataset
NewCustomerList.shape
```

```
Out[39]: (1000, 18)
```

```
In [40]: #Checking for null values
NewCustomerList.isnull().sum()
```

```
Out[40]: first_name      0
last_name    29
gender       0
past_3_years_bike_related_purchases  0
DOB         17
job_title    106
job_industry_category  165
wealth_segment  0
deceased_indicator  0
owns_car     0
tenure       0
address      0
postcode     0
state        0
country      0
property_valuation  0
Rank         0
Value        0
dtype: int64
```

There are missing values in 4 columns. They can be dropped or treated according to the nature of analysis

```
In [41]: #Checking for duplicate values
NewCustomerList.duplicated().sum()
```

```
Out[41]: 0
```

There are no duplicate values.

```
In [43]: #Checking for uniqueness of each column
NewCustomerList.nunique()
```

```
Out[43]: first_name      940
last_name    961
gender       3
past_3_years_bike_related_purchases  100
DOB         958
job_title    184
job_industry_category  9
wealth_segment  3
deceased_indicator  1
owns_car     2
tenure       23
address      1000
postcode     522
state        3
country      1
property_valuation  12
Rank         324
```

Value
dtype: int64

324

Exploring the columns

In [44]: `NewCustomerList.columns`

Out[44]: `Index(['first_name', 'last_name', 'gender',
'past_3_years_bike_related_purchases', 'DOB', 'job_title',
'job_industry_category', 'wealth_segment', 'deceased_indicator',
'owns_car', 'tenure', 'address', 'postcode', 'state', 'country',
'property_valuation', 'Rank', 'Value'],
dtype='object')`

In [45]: `NewCustomerList['gender'].value_counts()`

Out[45]:

Female	513
Male	470
U	17

Name: gender, dtype: int64

There are 17 columns with unknown/unspecified gender.

In [47]: `NewCustomerList['DOB'].value_counts()`

Out[47]:

1998-02-05	2
1978-01-15	2
1977-11-08	2
1951-11-28	2
1979-07-28	2
..	
1945-08-08	1
1943-08-27	1
1999-10-24	1
1976-01-24	1
1955-10-02	1

Name: DOB, Length: 958, dtype: int64

In [48]: `NewCustomerList['job_industry_category'].value_counts()`

Out[48]:

Financial Services	203
Manufacturing	199
Health	152
Retail	78
Property	64
IT	51
Entertainment	37
Argiculture	26
Telecommunications	25

Name: job_industry_category, dtype: int64

In [49]: `NewCustomerList['wealth_segment'].value_counts()`

Out[49]:

Mass Customer	508
High Net Worth	251
Affluent Customer	241

Name: wealth_segment, dtype: int64

In [50]:

NewCustomerList['state'].value_counts()

Out[50]:

NSW 506
VIC 266
QLD 228
Name: state, dtype: int64

In [51]:

NewCustomerList['owns_car'].value_counts()

Out[51]:

No 507
Yes 493
Name: owns_car, dtype: int64

In [52]:

NewCustomerList['deceased_indicator'].value_counts()

Out[52]:

N 1000
Name: deceased_indicator, dtype: int64

Exploring Customer Demographic Data Set

In [53]:

CustomerDemographic.head()

Out[53]:

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job
0	1	Laraine	Medendorp	F	93	1953-10-12	Executive Secretary
1	2	Eli	Bockman	Male	81	1980-12-16	Administrative Clerk
2	3	Arlin	Dearle	Male	61	1954-01-20	Recreation Manager
3	4	Talbot	NaN	Male	33	1961-10-03	
4	5	Sheila-kathryn	Calton	Female	56	1977-05-13	Senior Manager

In [54]:

CustomerDemographic.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 13 columns):
Column Non-Null Count Dtype
--- -
0 customer_id 4000 non-null int64
1 first_name 4000 non-null object
2 last_name 3875 non-null object
3 gender 4000 non-null object
4 past_3_years_bike_related_purchases 4000 non-null int64
5 DOB 3913 non-null datetime64[ns]
6 job_title 3494 non-null object
7 job_industry_category 3344 non-null object
8 wealth_segment 4000 non-null object
9 deceased_indicator 4000 non-null object
10 default 3698 non-null object

```

11  owns_car                4000 non-null    object
12  tenure                  3913 non-null  float64
dtypes: datetime64[ns](1), float64(1), int64(2), object(9)
memory usage: 406.4+ KB

```

```

In [55]: #Checking for null values
         CustomerDemographic.isnull().sum()

```

```

Out[55]: customer_id                0
         first_name                 0
         last_name                125
         gender                    0
         past_3_years_bike_related_purchases  0
         DOB                      87
         job_title                 506
         job_industry_category     656
         wealth_segment            0
         deceased_indicator         0
         default                  302
         owns_car                  0
         tenure                    87
         dtype: int64

```

There are missing values in 5 columns. They can be dropped or treated according to the nature of analysis

```

In [56]: #Checking for duplicate data
         CustomerDemographic.duplicated().sum()

```

```

Out[56]: 0

```

There are no duplicate values.

```

In [57]: #Checking for uniqueness of each column
         CustomerDemographic.nunique()

```

```

Out[57]: customer_id                4000
         first_name                 3139
         last_name                 3725
         gender                     6
         past_3_years_bike_related_purchases  100
         DOB                      3448
         job_title                 195
         job_industry_category      9
         wealth_segment             3
         deceased_indicator          2
         default                   90
         owns_car                   2
         tenure                     22
         dtype: int64

```

Exploring the columns

```

In [70]: CustomerDemographic.columns

```

```

Out[70]: Index(['customer_id', 'first_name', 'last_name', 'gender',
               'past_3_years_bike_related_purchases', 'DOB', 'job_title',
               'job_industry_category', 'wealth_segment', 'deceased_indicator',

```

```
'default', 'owns_car', 'tenure'],
dtype='object')
```

```
In [71]: CustomerDemographic['gender'].value_counts()
```

```
Out[71]: Female      2039
Male        1873
Unspecific    88
Name: gender, dtype: int64
```

Certain categories are not correctly titled. The names in these categories are re-named.

```
In [72]: #Re-naming the categories
CustomerDemographic['gender'] = CustomerDemographic['gender'].replace('F', 'Female').
```

```
In [73]: CustomerDemographic['gender'].value_counts()
```

```
Out[73]: Female      2039
Male        1873
Unspecific    88
Name: gender, dtype: int64
```

```
In [74]: CustomerDemographic['past_3_years_bike_related_purchases'].value_counts()
```

```
Out[74]: 16    56
19    56
67    54
20    54
2     50
..
8     28
95    27
85    27
86    27
92    24
Name: past_3_years_bike_related_purchases, Length: 100, dtype: int64
```

```
In [75]: CustomerDemographic['DOB'].value_counts()
```

```
Out[75]: 1978-01-30    7
1964-07-08    4
1962-12-17    4
1978-08-19    4
1977-05-13    4
..
1989-06-16    1
1998-09-30    1
1985-03-11    1
1989-10-23    1
1991-11-05    1
Name: DOB, Length: 3448, dtype: int64
```

```
In [76]: CustomerDemographic['job_title'].value_counts()
```

```
Out[76]: Business Systems Development Analyst    45
Tax Accountant                                44
Social Worker                                  44
Internal Auditor                              42
```

```

Recruiting Manager      41
..
Database Administrator I  4
Health Coach I          3
Health Coach III        3
Research Assistant III   3
Developer I             1
Name: job_title, Length: 195, dtype: int64

```

```
In [77]: CustomerDemographic['job_industry_category'].value_counts()
```

```

Out[77]: Manufacturing      799
Financial Services    774
Health                602
Retail                358
Property              267
IT                    223
Entertainment         136
Agriculture           113
Telecommunications    72
Name: job_industry_category, dtype: int64

```

```
In [78]: CustomerDemographic['wealth_segment'].value_counts()
```

```

Out[78]: Mass Customer      2000
High Net Worth      1021
Affluent Customer    979
Name: wealth_segment, dtype: int64

```

```
In [79]: CustomerDemographic['deceased_indicator'].value_counts()
```

```

Out[79]: N      3998
Y         2
Name: deceased_indicator, dtype: int64

```

```
In [80]: CustomerDemographic['default'].value_counts()
```

```

Out[80]: 100      113
1          112
-1         111
-100       99
Û;Û$Û£     53
...
testâ testâ«      31
/dev/null; touch /tmp/blns.fail ; echo      30
âââtestââ      29
ì,ëë°í ë¥´      27
,ãã»:ã»ãâ( â» Ĭ â» )ãã»:ã»ãâ      25
Name: default, Length: 90, dtype: int64

```

```
In [81]: CustomerDemographic = CustomerDemographic.drop('default', axis=1)
```

The values are inconsistent, hence dropping the column.

```
In [82]: CustomerDemographic.head(5)
```

```
Out[82]: customer_id  first_name  last_name  gender  past_3_years_bike_related_purchases  DOB  job
```

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job
0	1	Laraine	Medendorp	Female		93 1953-10-12	Exec Sec
1	2	Eli	Bockman	Male		81 1980-12-16	Administ C
2	3	Arlin	Dearle	Male		61 1954-01-20	Recr Ma
3	4	Talbot	NaN	Male		33 1961-10-03	
4	5	Sheila-kathryn	Calton	Female		56 1977-05-13	Senior I

In [83]:

CustomerDemographic['owns_car'].value_counts()

Out[83]:

Yes 2024
No 1976
Name: owns_car, dtype: int64

In [84]:

CustomerDemographic['tenure'].value_counts()

Out[84]:

7.0 235
5.0 228
11.0 221
10.0 218
16.0 215
8.0 211
18.0 208
12.0 202
9.0 200
14.0 200
6.0 192
13.0 191
4.0 191
17.0 182
15.0 179
1.0 166
3.0 160
19.0 159
2.0 150
20.0 96
22.0 55
21.0 54
Name: tenure, dtype: int64

Exploring Customer Address Data Set

In [85]:

CustomerAddress.head(5)

Out[85]:

	customer_id	address	postcode	state	country	property_valuation
0	1	060 Morning Avenue	2016	New South Wales	Australia	10
1	2	6 Meadow Vale Court	2153	New South Wales	Australia	10
2	4	0 Holy Cross Court	4211	QLD	Australia	9

	customer_id	address	postcode	state	country	property_valuation
3	5	17979 Del Mar Point	2448	New South Wales	Australia	4
4	6	9 Oakridge Court	3216	VIC	Australia	9

In [86]:

```
CustomerAddress.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3999 entries, 0 to 3998
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customer_id           3999 non-null   int64
1   address                3999 non-null   object
2   postcode               3999 non-null   int64
3   state                  3999 non-null   object
4   country                3999 non-null   object
5   property_valuation     3999 non-null   int64
dtypes: int64(3), object(3)
memory usage: 187.6+ KB
```

In [87]:

```
#Checking for null values.
CustomerAddress.isnull().sum()
```

Out[87]:

```
customer_id      0
address          0
postcode         0
state            0
country          0
property_valuation 0
dtype: int64
```

There are no null values.

In [88]:

```
#Checking for duplicate values
CustomerAddress.duplicated().sum()
```

Out[88]:

0

There are no duplicate values.

In [89]:

```
#Checking for uniqueness of each column
CustomerAddress.nunique()
```

Out[89]:

```
customer_id      3999
address          3996
postcode         873
state            5
country          1
property_valuation 12
dtype: int64
```

Exploring the columns

In [90]:

```
CustomerAddress['postcode'].value_counts()
```

```
Out[90]: 2170    31
         2155    30
         2145    30
         2153    29
         3977    26
         ..
         3808     1
         3114     1
         4721     1
         4799     1
         3089     1
         Name: postcode, Length: 873, dtype: int64
```

```
In [91]: CustomerAddress['state'].value_counts()
```

```
Out[91]: NSW          2054
         VIC           939
         QLD           838
         New South Wales  86
         Victoria       82
         Name: state, dtype: int64
```

```
In [92]: CustomerAddress['country'].value_counts()
```

```
Out[92]: Australia    3999
         Name: country, dtype: int64
```

```
In [93]: CustomerAddress['property_valuation'].value_counts()
```

```
Out[93]: 9      647
         8      646
         10     577
         7      493
         11     281
         6      238
         5      225
         4      214
         12     195
         3      186
         1      154
         2      143
         Name: property_valuation, dtype: int64
```

All the columns appear to have consistent and correct information.